


```
from google.colab import drive  
drive.mount('/content/gdrive')
```

Mounted at /content/gdrive

Capstone Project: Predicting Hospital Readmission Risk in Diabetic Patients



Double-click (or enter) to edit

Import Libraries

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import LabelEncoder, StandardScaler  
from sklearn.impute import SimpleImputer  
from sklearn.linear_model import LogisticRegression  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.svm import SVC  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.metrics import accuracy_score, confusion_matrix, classificat.
```

Load and Explore Data

```
data = pd.read_csv("/content/gdrive/MyDrive/diabetic_data.csv")
```

```

# Display basic info
print("Dataset Shape:", data.shape)
print("\nColumns:\n", data.columns.tolist())
print("\nPreview:\n", data.head())

```

Dataset Shape: (17469, 50)

Columns:

```

['encounter_id', 'patient_nbr', 'race', 'gender', 'age', 'weight', 'admis

```

Preview:

	encounter_id	patient_nbr	race	gender	age	weight
0	2278392	8222157	Caucasian	Female	[0-10)	?
1	149190	55629189	Caucasian	Female	[10-20)	?
2	64410	86047875	AfricanAmerican	Female	[20-30)	?
3	500364	82442376	Caucasian	Male	[30-40)	?
4	16680	42519267	Caucasian	Male	[40-50)	?

	admission_type_id	discharge_disposition_id	admission_source_id
0	6	25	1
1	1	1	7
2	1	1	7
3	1	1	7
4	1	1	7

	time_in_hospital	... citoglipton	insulin	glyburide-metformin
0	1	...	No	No
1	3	...	No	Up
2	2	...	No	No
3	2	...	No	Up
4	1	...	No	Steady

	glipizide-metformin	glimepiride-pioglitazone	metformin-rosiglitazone
0	No	No	No
1	No	No	No
2	No	No	No
3	No	No	No
4	No	No	No

	metformin-pioglitazone	change	diabetesMed	readmitted
0	No	No	No	NO
1	No	Ch	Yes	>30
2	No	No	Yes	NO
3	No	Ch	Yes	NO
4	No	Ch	Yes	NO

[5 rows x 50 columns]

Data cleaning

```

# Replace '?' with NaN
data.replace('?', np.nan, inplace=True)

```

```
# Drop columns that are not useful for prediction
data.drop(['encounter_id', 'patient_nbr', 'weight', 'payer_code', 'medica

# Check how many missing values remain
print("\nMissing values per column:\n", data.isnull().sum())
```

Missing values per column:

race	291
gender	0
age	0
admission_type_id	0
discharge_disposition_id	0
admission_source_id	0
time_in_hospital	0
num_lab_procedures	0
num_procedures	0
num_medications	0
number_outpatient	0
number_emergency	0
number_inpatient	0
diag_1	7
diag_2	104
diag_3	498
number_diagnoses	0
max_glu_serum	15754
A1Cresult	14213
metformin	0
repaglinide	0
nateglinide	0
chlorpropamide	0
glimepiride	0
acetohexamide	0
glipizide	0
glyburide	0
tolbutamide	0
pioglitazone	0
rosiglitazone	0
acarbose	0
miglitol	0
troglitazone	0
tolazamide	0
examide	0
citoglipton	0
insulin	0
glyburide-metformin	0
glipizide-metformin	0
glimepiride-pioglitazone	0
metformin-rosiglitazone	0
metformin-pioglitazone	0
change	0
diabetesMed	0
readmitted	0
dtype: int64	

/tmp/ipython-input-1993658940.py:2: FutureWarning: Downcasting behavior in

```
data.replace('?', np.nan, inplace=True)
```

Feature Engineering

```
# Convert target variable 'readmitted' into binary
# '<30' and '>30' = 1 (readmitted), 'NO' = 0 (not readmitted)
data['readmitted'] = data['readmitted'].replace({'NO': 0, '<30': 1, '>30': 1})

# Encode all categorical (string) columns
label_encoder = LabelEncoder()
categorical_cols = data.select_dtypes(include=['object']).columns

for col in categorical_cols:
    data[col] = label_encoder.fit_transform(data[col].astype(str))

print("Categorical columns encoded:", list(categorical_cols))
```

```
/tmp/ipython-input-136016837.py:3: FutureWarning: Downcasting behavior in  
data['readmitted'] = data['readmitted'].replace({'NO': 0, '<30': 1, '>30
```

Handle Missing Values and Scale Data

```
# Fill missing values with the most frequent value  
imputer = SimpleImputer(strategy='most_frequent')  
data_imputed = pd.DataFrame(imputer.fit_transform(data), columns=data.columns)  
  
# Scale numeric features  
scaler = StandardScaler()  
X = data_imputed.drop('readmitted', axis=1)  
y = data_imputed['readmitted']  
X_scaled = scaler.fit_transform(X)  
  
print("After preprocessing, data shape:", X_scaled.shape)
```

```
After preprocessing, data shape: (17469, 44)
```

Train-Test Split

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)  
  
print("Training set size:", X_train.shape)  
print("Testing set size:", X_test.shape)
```

```
Training set size: (13975, 44)  
Testing set size: (3494, 44)
```

Feature Scaling


```
scaler = StandardScaler()  
X_train = scaler.fit_transform(X_train)  
X_test = scaler.transform(X_test)
```

```
import matplotlib.pyplot as plt  
import seaborn as sns  
  
num_cols = ['time_in_hospital', 'num_lab_procedures', 'num_medications',  
            'num_surgeries', 'num_physicians', 'num_diagnoses']  
  
plt.figure(figsize=(12,8))
```

```

for i, col in enumerate(num_cols, 1):
    plt.subplot(2, 3, i)
    sns.histplot(data[col], bins=30, kde=True)
    plt.title(f'Distribution of {col}')
plt.tight_layout()
plt.show()

```



```


plt.figure(figsize=(6,5))
sns.countplot(x='readmitted', data=data, palette='Set2')
plt.title('Distribution of Readmission Categories')
plt.xlabel('Readmission Status')
plt.ylabel('Count')
plt.show()

```


```
/tmp/ipython-input-1917754014.py:2: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed
```

```
sns.countplot(x='readmitted', data=data, palette='Set2')
```




```
plt.figure(figsize=(10,8))
sns.heatmap(data[num_cols].corr(), annot=True, cmap='Blues')
plt.title('Correlation Between Numerical Features')
plt.show()
```



```
plt.figure(figsize=(8,5))
sns.barplot(x='age', y='time_in_hospital', data=data, palette='magma')
plt.title('Average Hospital Stay Duration by Age Group')
plt.xticks(rotation=45)
plt.show()
```

```
/tmp/ipython-input-4067393693.py:2: FutureWarning:  
  Passing `palette` without assigning `hue` is deprecated and will be removed  
  sns.barplot(x='age', y='time_in_hospital', data=data, palette='magma')  
  
  Average Hospital Stay Duration by Age Group
```



Logistic Regression


```
log_model = LogisticRegression(max_iter=2000)  
log_model.fit(X_train, y_train)  
  
y_pred_log = log_model.predict(X_test)  
  
print("Logistic Regression Accuracy:", accuracy_score(y_test, y_pred_log))  
print("\nClassification Report:\n", classification_report(y_test, y_pred_log))  
  
# Confusion Matrix  
plt.figure(figsize=(5,4))  
sns.heatmap(confusion_matrix(y_test, y_pred_log), annot=True, fmt='d', cmap='Blues')  
plt.title("Logistic Regression - Confusion Matrix")  
plt.xlabel("Predicted")  
plt.ylabel("Actual")  
plt.show()
```

Logistic Regression Accuracy: 0.6210646823125358

Classification Report:

	precision	recall	f1-score	support
0	0.61	0.73	0.67	1811
1	0.64	0.50	0.56	1683
accuracy			0.62	3494
macro avg	0.62	0.62	0.61	3494
weighted avg	0.62	0.62	0.62	3494

Logistic Regression - Confusion Matrix



Random Forest Classifier

```
rf_model = RandomForestClassifier(n_estimators=200, random_state=42)
rf_model.fit(X_train, y_train)

y_pred_rf = rf_model.predict(X_test)

print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))
print("\nClassification Report:\n", classification_report(y_test, y_pred_rf))


plt.figure(figsize=(5,4))
sns.heatmap(confusion_matrix(y_test, y_pred_rf), annot=True, fmt='d', cmap='Blues')
plt.title("Random Forest - Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

Random Forest Accuracy: 0.6313680595306239

Classification Report:

	precision	recall	f1-score	support
0	0.63	0.69	0.66	1811
1	0.63	0.57	0.60	1683
accuracy			0.63	3494
macro avg	0.63	0.63	0.63	3494
weighted avg	0.63	0.63	0.63	3494

Random Forest - Confusion Matrix



K-Nearest Neighbors

```
knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(X_train, y_train)

y_pred_knn = knn_model.predict(X_test)


print("KNN Accuracy:", accuracy_score(y_test, y_pred_knn))
print("\nClassification Report:\n", classification_report(y_test, y_pred_knn))

plt.figure(figsize=(5,4))
sns.heatmap(confusion_matrix(y_test, y_pred_knn), annot=True, fmt='d', cmap='viridis')
plt.title("KNN - Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

KNN Accuracy: 0.5729822552947911

Classification Report:

	precision	recall	f1-score	support
0	0.58	0.62	0.60	1811
1	0.56	0.52	0.54	1683
accuracy			0.57	3494
macro avg	0.57	0.57	0.57	3494
weighted avg	0.57	0.57	0.57	3494



Compare All Models

```
model_results = {
    "Model": ["Logistic Regression", "Random Forest", "KNN"],
    "Accuracy": [
        accuracy_score(y_test, y_pred_log),
        accuracy_score(y_test, y_pred_rf),
        accuracy_score(y_test, y_pred_knn)
    ]
}
results_df = pd.DataFrame(model_results)
print("\n\n FINAL MODEL PERFORMANCE")
print(results_df)

plt.figure(figsize=(8,4))
sns.barplot(x="Model", y="Accuracy", data=results_df, palette="viridis")
plt.title("Comparison of Model Accuracies")
```

```
plt.xticks(rotation=30)
plt.show()
```


```
FINAL MODEL PERFORMANCE
      Model  Accuracy
0  Logistic Regression  0.621065
1          Random Forest  0.631368
2              KNN  0.572982
```

```
/tmp/ipython-input-268332014.py:14: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed
```

```
sns.barplot(x="Model", y="Accuracy", data=results_df, palette="viridis")
```

Comparison of Model Accuracies



```
print("\nFINAL CONCLUSION:")
print("Random Forest performed the best and is recommended for predicting
```

FINAL CONCLUSION:

Random Forest performed the best and is recommended for predicting readmis

```
import numpy as np
```

```
feature_importances = rf_model.feature_importances_
```